

# 使用模拟集群运行分布式tensorflow

Created by zhaomingxing, last modified on Feb 10, 2017

第一步：创建多个docker container，我这里创建了2个参数服务器，ps0和ps1；2个worker，w0和w1。参数服务器ps和worker可以设置一个或多个。使用的是docker官方的镜像，最新版的tensorflow/tensorflow的镜像。

```
docker run -d -v ~/Desktop/zmx/dockers:/tf/mnist --name ps0 tensorflow/tensorflow
```

```
docker run -d -v ~/Desktop/zmx/dockers:/tf/mnist --name ps1 tensorflow/tensorflow
```

```
docker run -d -v ~/Desktop/zmx/dockers:/tf/mnist --name w0 tensorflow/tensorflow
```

```
docker run -d -v ~/Desktop/zmx/dockers:/tf/mnist --name w1 tensorflow/tensorflow
```

这里的~/Desktop/zmx/dockers是本地存放分布式mnist示例代码（distribute\_version\_MNIST.py）的目录，也可以不进行目录映射，但是要把同一份代码分别拷贝到每个docker container中，因为后面每个docker container都要运行这份代码，只是运行命令稍有区别。执行完上述命令之后查看已经生成的docker container如下：

| CONTAINER ID | IMAGE                 | COMMAND           | CREATED      | STATUS                    | PORTS              | NAMES |
|--------------|-----------------------|-------------------|--------------|---------------------------|--------------------|-------|
| 2a4398ac72e5 | schickling/rust       | "bash"            | 16 hours ago | Exited (137) 10 hours ago |                    | gdb   |
| 25b2c018ffcc | tensorflow/tensorflow | "/run_jupyter.sh" | 43 hours ago | Up 20 minutes             | 6006/tcp, 8888/tcp | w1    |
| 11b95874fccf | tensorflow/tensorflow | "/run_jupyter.sh" | 43 hours ago | Up 12 seconds             | 6006/tcp, 8888/tcp | w0    |
| cadf4e75d491 | tensorflow/tensorflow | "/run_jupyter.sh" | 43 hours ago | Up 6 seconds              | 6006/tcp, 8888/tcp | ps1   |
| 553a13f9b79f | tensorflow/tensorflow | "/run_jupyter.sh" | 43 hours ago | Up 8 seconds              | 6006/tcp, 8888/tcp | ps0   |

第二步：使用ifconfig命令查看上述的四个docker container的ip地址。

ps0: 172.17.0.4

ps1: 172.17.0.5

w0: 172.17.0.3

w1: 172.17.0.2


以ps0运行ifconfig的结果为例：

```
zhaomingxing@zhaomingxingdeMacBook-Pro:~$ docker exec -it ps0 /bin/bash
root@553a13f9b79f:/notebooks# ifconfig
eth0      Link encap:Ethernet  HWaddr 02:42:ac:11:00:04
          inet addr:172.17.0.4  Bcast:0.0.0.0  Mask:255.255.0.0
          inet6 addr: fe80::42:acff:fe11:4/64 Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:66 errors:0 dropped:0 overruns:0 frame:0
          TX packets:13 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:3580 (3.5 KB)  TX bytes:998 (998.0 B)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
          UP LOOPBACK RUNNING  MTU:65536  Metric:1
          RX packets:0 errors:0 dropped:0 overruns:0 frame:0
          TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1
          RX bytes:0 (0.0 B)  TX bytes:0 (0.0 B)

root@553a13f9b79f:/notebooks#
```

第三步：使用docker exec -ti ps1(docker container的name或ID) /bin/bash分别打开每个docker container的终端，切换到分布式MNIST代码所在的目录。



第四步：在上述的四个终端中分别运行分布式MNIST的代码distribute\_version\_MNIST.py：

ps0执行：

```
python distribute_version_MNIST.py --ps_hosts=172.17.0.4:2222,172.17.0.5:2222 --worker_hosts=172.17.0.3:2222,172.17.0.2:2222 --job_name=ps --task_index=0
```

ps1执行：

```
python distribute_version_MNIST.py --ps_hosts=172.17.0.4:2222,172.17.0.5:2222 --worker_hosts=172.17.0.3:2222,172.17.0.2:2222 --job_name=ps --task_index=1
```

w0执行：

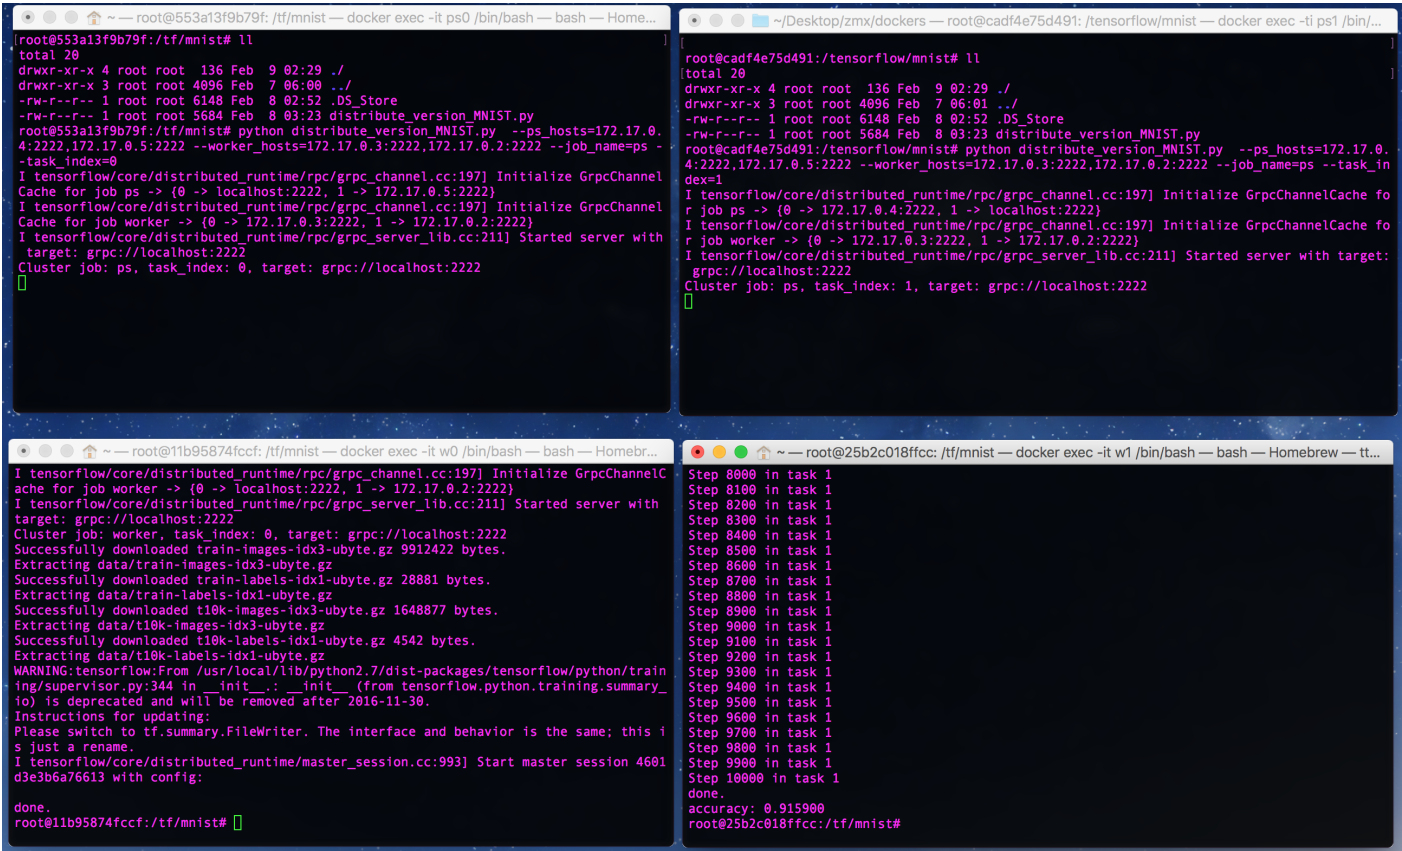
```
python distribute_version_MNIST.py --ps_hosts=172.17.0.4:2222,172.17.0.5:2222 --worker_hosts=172.17.0.3:2222,172.17.0.2:2222 --job_name=worker --task_index=0
```

w1执行：

```
python distribute_version_MNIST.py --ps_hosts=172.17.0.4:2222,172.17.0.5:2222 --worker_hosts=172.17.0.3:2222,172.17.0.2:2222 --job_name=worker --task_index=1
```

**注意：**w1执行几秒后会退出一次，再执行一次相同的命令即可。上述端口号(2222)是可以随便指定的，只要是docker container上的空闲端口号即可。


运行结果如下：



从w1的输出结果来看，运行10000步之后，测试集上的准确率达到了0.915900。

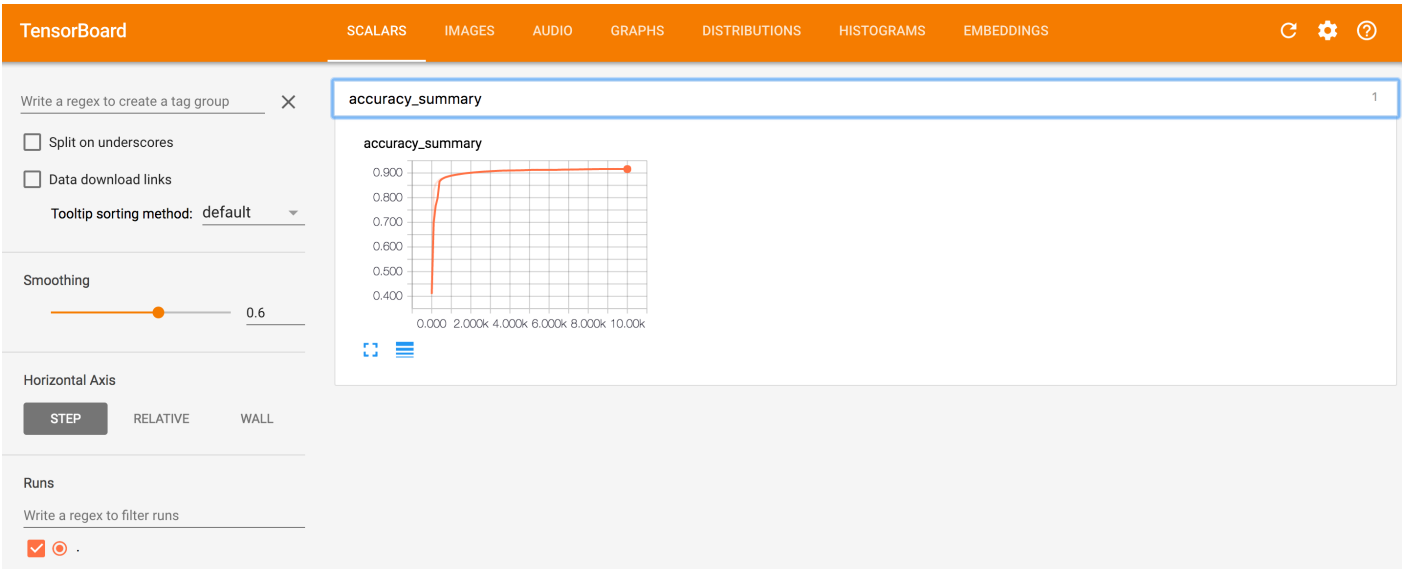
第五步：使用tensorboard观察数据流图和程序执行过程中的accuracy的变化情况。

tensorboard的执行如下：



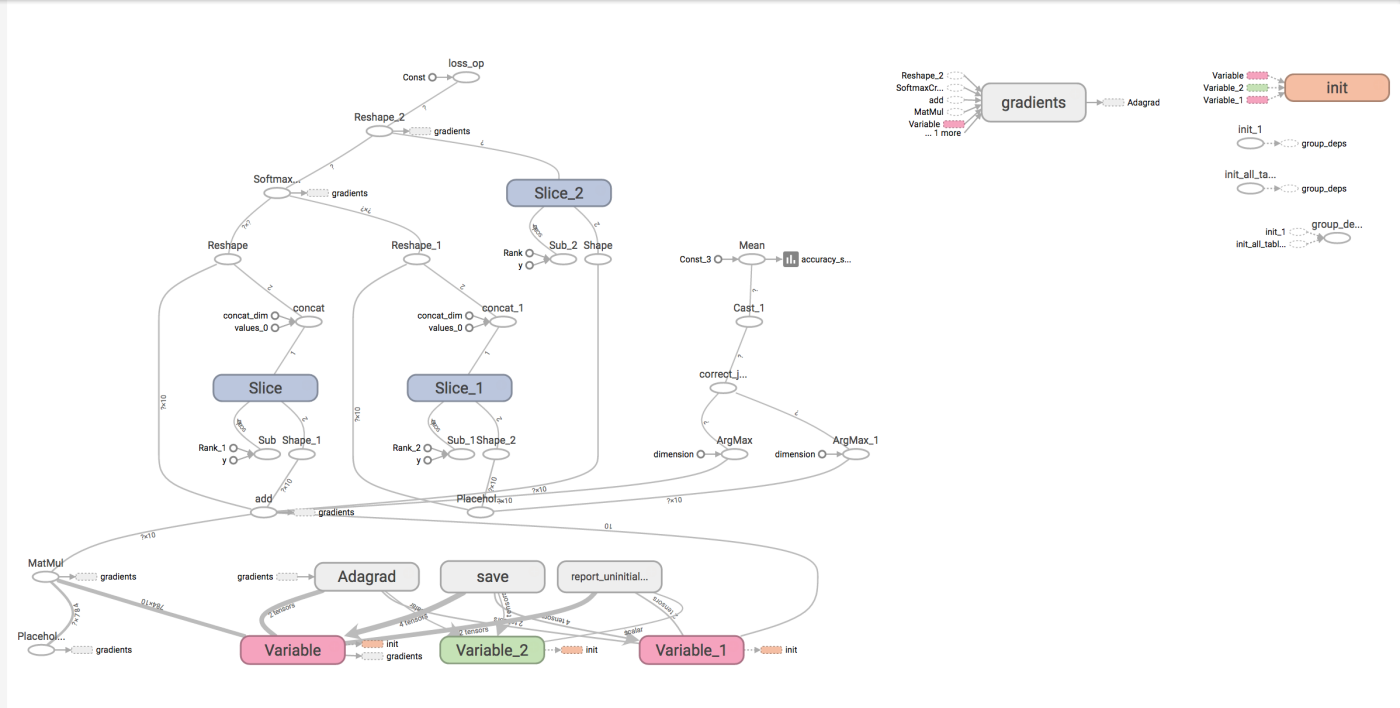
按照图片中的指示：在浏览器地址栏中输入<http://192.168.0.62:6006>即可。

选择“SCALARS”栏，点击accuracy\_summary可见下图：



上面的折线图就是训练过程中，测试集上的预测准确率的变化。

选择GRAPHS栏，可以看到下图：



最后附上分布式MNIST的源码：

```
distribute_version_MNIST.py


1  #!/usr/bin/python
2  #-*- coding:utf-8 -*-
3  """
4
5  =====
6  author: 赵明星
7  desc:   分布式tensorflow实现手写数字识别MNIST。
8  =====
9  """
10 import sys
11 reload(sys)
12 sys.setdefaultencoding('utf-8')
13
14 import tensorflow as tf
15 from tensorflow.examples.tutorials.mnist import input_data
16 import os
17
18 tf.app.flags.DEFINE_string("ps_hosts",
19                             "",
20                             "Comma-separated list of hostname:port pairs")
21 tf.app.flags.DEFINE_string("worker_hosts",
22                             "",
23                             "Comma-separated list of hostname:port pairs")
24 tf.app.flags.DEFINE_string("job_name", "", "One of 'ps', 'worker'")
25 tf.app.flags.DEFINE_integer("task_index", 0, "Index of task within the job")
26 FLAGS = tf.app.flags.FLAGS
27
28 def main(_):
29     ps_hosts = FLAGS.ps_hosts.split(",")
30     worker_hosts = FLAGS.worker_hosts.split(",")
31     # Create a cluster from the parameter server and worker hosts.
32     cluster = tf.train.ClusterSpec({"ps": ps_hosts, "worker": worker_hosts})
33     # Create and start a server for the local task.
34     server = tf.train.Server(cluster,
35                             job_name=FLAGS.job_name,
36                             task_index=FLAGS.task_index)
37     print("Cluster job: %s, task_index: %d, target: %s" % (FLAGS.job_name,
38                                                             FLAGS.task_index,
39                                                             server.target))
40
41     if FLAGS.job_name == "ps":
```

```

39     server.join()
40 elif FLAGS.job_name == "worker":
41     # Assigns ops to the local worker by default.
42     with tf.device(tf.train.replica_device_setter(
43         worker_device="/job:worker/task:%d" % FLAGS.task_index,
44         cluster=cluster)):
45         # Build model ...
46         mnist = input_data.read_data_sets("data", one_hot=True)
47
48         # Create the model
49         x = tf.placeholder(tf.float32, [None, 784])
50         W = tf.Variable(tf.zeros([784, 10]))
51         b = tf.Variable(tf.zeros([10]))
52         y = tf.matmul(x, W) + b
53         # Define loss and optimizer
54         y_ = tf.placeholder(tf.float32, [None, 10])
55         cross_entropy = tf.reduce_mean(
56             tf.nn.softmax_cross_entropy_with_logits(y, y_),
57             name="loss_op")
58         # loss_summary = tf.summary.scalar("loss", cross_entropy)
59         global_step = tf.Variable(0)
60         train_op = tf.train.AdagradOptimizer(0.01).minimize(
61             cross_entropy, global_step=global_step)
62         # Test trained model
63         correct_prediction = tf.equal(tf.argmax(y, 1),
64                                     tf.argmax(y_, 1),
65                                     name="correct_judge_op")
66         accuracy = tf.reduce_mean(tf.cast(correct_prediction,
67                                     tf.float32))
68         accuracy_summary = tf.summary.scalar("accuracy_summary",
69                                     accuracy)
70
71         saver = tf.train.Saver()
72         summary_op = tf.summary.merge_all()
73         init_op = tf.global_variables_initializer()
74     if not os.path.exists("mnist_log"):
75         os.mkdir("mnist_log")
76     # Create a "Supervisor", which oversees the training process.
77     sv = tf.train.Supervisor(is_chief=(FLAGS.task_index == 0),
78                             logdir="mnist_log",
79                             init_op=init_op,
80                             summary_op=summary_op,
81                             saver = saver,
82                             global_step=global_step,
83                             save_model_secs=600)
84     # The supervisor takes care of session initialization
85     # and restoring from a checkpoint.
86     sess = sv.prepare_or_wait_for_session(server.target)
87     writer = tf.summary.FileWriter("mnist_log", sess.graph)
88     # Start queue runners for the input pipelines (if any).
89     sv.start_queue_runners(sess)
90     # Loop until the supervisor shuts down (or 2000 steps have completed).
91     step = 0
92     while not sv.should_stop() and step < 10000:
93         batch_xs, batch_ys = mnist.train.next_batch(100)
94         _, step = sess.run([train_op, global_step],
95                             feed_dict={x: batch_xs, y_: batch_ys})
96         if step % 100 == 0 and FLAGS.task_index != 0:
97             res = sess.run(summary_op,
98                             feed_dict={x: mnist.test.images,
99                                     y_: mnist.test.labels})
100             writer.add_summary(res, step)
101             print("Step {0} in task {1}".format(step, FLAGS.task_index))
102     print("done.")
103     if not os.path.exists("variables_saved"):
104         os.mkdir("variables_saved")
105     saver.save(sess, "variables_saved/variables")

```

```
106         if FLAGS.task_index != 0:
107             print("accuracy: %f" % sess.run(accuracy,
108                                             feed_dict={x: mnist.test.images,
109                                                     y: mnist.test.labels}))
110     if __name__ == "__main__":
111         tf.app.run()
```

 [Like](#) Be the first to like this

No labels